

# Model\_v1

July 14, 2025

## 1 Model V1

```
[37]: best_config = {'lstm_bidirectional': False,
                    'batch_size': 64,
                    'lstm_hidden': 96,
                    'lstm_layers': 2,
                    'lstm_dropout': 0.1,
                    'gat_hidden': 96,
                    'gat_dropout': 0.1,
                    'gat_alpha': 0.1,
                    'final_dropout': 0.2,
                    'learning_rate': 0.0007023432254879037,
                    'lstm_weight_decay': 0.0005708066170581489,
                    'gat_weight_decay': 1.6797946539308408e-05,
                    'final_weight_decay': 0.00017357100923169329}
```

```
[38]: import os
import torch
import torch.nn as nn
import torch.nn.functional as F
import pandas as pd
import numpy as np
import random
from sklearn.preprocessing import StandardScaler
from tqdm import tqdm

def set_seed(seed=42):
    torch.manual_seed(seed)
    np.random.seed(seed)
    random.seed(seed)
    if torch.cuda.is_available():
        torch.cuda.manual_seed(seed)
        torch.cuda.manual_seed_all(seed)
        torch.backends.cudnn.deterministic = True
        torch.backends.cudnn.benchmark = False
    os.environ['PYTHONHASHSEED'] = str(seed)
set_seed(42)
```

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Using device:", device)
```

Using device: cpu

## 2 Load Data and Split Training and Testing

```
[39]: import os
import pandas as pd
import numpy as np

tickers = ['AAPL', 'AMZN', 'BA', 'COST', 'JNJ', 'NVDA', 'TMO', 'TSLA', 'VLO']
data_dir = "Data"
features = ['close', 'volume', 'log_return']

all_data = {}
min_length = float('inf')

for stock in tickers:
    df = pd.read_csv(os.path.join(data_dir,
    ↪f"{stock}_with_sentiment_features_with_product.csv"))

    # If log_return is not already present, add it
    if 'log_return' not in df.columns:
        df['log_return'] = np.log(df['close'] / df['close'].shift(1))

    # Fill missing price features using forward fill (safer for price data)
    price_feats = ['open', 'high', 'low', 'close', 'volume',
                    'ann_return_1w', 'ann_return_2w', 'ann_return_1m',
                    'rolling_vol_7d', 'ann_volatility', 'macd_1w_1m',
    ↪'log_return']
    for feat in price_feats:
        if feat in df.columns:
            df[feat] = df[feat].ffill()

    # Fill missing sentiment/news features with 0 (standard for no news days)
    for feat in ['news_count', 'mean_sentiment', 'sentiment_variance',
    ↪'product']:
        if feat in df.columns:
            df[feat] = df[feat].fillna(0)

    # Drop rows where the main features are still missing (e.g., very early
    ↪rows)
    df = df.dropna(subset=features).reset_index(drop=True)
    all_data[stock] = df
```

```

min_length = min(min_length, len(df))

# Align all stocks to same length from the end (for parallel modeling)
for stock in tickers:
    all_data[stock] = all_data[stock].tail(min_length).reset_index(drop=True)

# Unified date array
dates = all_data[tickers[0]]['date'].values
total_len = len(dates)

# Train/test split
test_size = 0.3
split_idx = int(total_len * (1 - test_size))
train_dates = dates[:split_idx]
test_dates = dates[split_idx:]

print(f"Train: {train_dates[0]}~{train_dates[-1]}")
print(f"Test: {test_dates[0]}~{test_dates[-1]}")

```

Train: 2021-01-05~2024-02-01  
Test: 2024-02-02~2025-05-30

```

[40]: scalers = {}
for stock in tickers:
    train_df = all_data[stock][all_data[stock]['date'].isin(train_dates)]
    scaler = StandardScaler()
    scaler.fit(train_df[features])
    scalers[stock] = scaler
print("Scaler finished.")

```

Scaler finished.

### 3 Build Training Correlation Matrix

```

[41]: train_price_data = pd.DataFrame({'date': train_dates})
for stock in tickers:
    train_price_data[stock] = all_data[stock][all_data[stock]['date'].
isin(train_dates)]['close'].values
train_corr_matrix = train_price_data.drop(columns='date').corr()

edge_index, edge_attr = [], []
for i in range(len(tickers)):
    for j in range(len(tickers)):
        if i != j:
            edge_index.append([i, j])
            edge_attr.append(train_corr_matrix.iloc[i, j])

```

```

edge_index = torch.tensor(edge_index, dtype=torch.long).t().contiguous()
edge_attr = torch.tensor(edge_attr, dtype=torch.float32)
print("Edges weighted prepared.")

```

Edges weighted prepared.

## 4 Define LSTM and GAT

```

[ ]: # LSTM
class StockLSTMEncoder(nn.Module):
    def __init__(self, input_dim, hidden_dim, num_layers=1, dropout=0.0,
        ↪bidirectional=False):
        super().__init__()
        self.lstm = nn.LSTM(input_dim, hidden_dim, batch_first=True)
    def forward(self, x):
        output, (h_n, _) = self.lstm(x)
        return output[:, -1, :]

# GAT
class WeightedGATConv(nn.Module):
    def __init__(self, in_channels, out_channels, dropout=0.0, alpha=0.2):
        super().__init__()
        self.lin = nn.Linear(in_channels, out_channels)
        self.dropout = nn.Dropout(dropout)
        self.leaky_relu = nn.LeakyReLU(alpha)

    def forward(self, x, edge_index, edge_weight):
        x = self.lin(x)
        x = self.dropout(x)
        num_nodes = x.size(0)
        agg = torch.zeros_like(x)
        for idx in range(edge_index.size(1)):
            src = edge_index[0, idx]
            tgt = edge_index[1, idx]
            agg[tgt] += edge_weight[idx] * x[src]
        return self.leaky_relu(agg + x)

class GATEncoder(nn.Module):
    def __init__(self, in_dim, hidden_dim, out_dim, dropout=0.0, alpha=0.2):
        super().__init__()
        self.gat1 = WeightedGATConv(in_dim, hidden_dim, dropout, alpha)
        self.gat2 = WeightedGATConv(hidden_dim, out_dim, dropout, alpha)
    def forward(self, x, edge_index, edge_weight):
        x = self.gat1(x, edge_index, edge_weight)
        x = self.gat2(x, edge_index, edge_weight)
        return x

```

```
[61]: # Initialize base on config
lstm_model = StockLSTMEncoder(
    input_dim=3,
    hidden_dim=best_config['lstm_hidden'],
    num_layers=best_config['lstm_layers'],
    dropout=best_config['lstm_dropout'],
    bidirectional=best_config['lstm_bidirectional']
).to(device)

gat_input_dim = best_config['lstm_hidden'] * (2 if
↳best_config['lstm_bidirectional'] else 1)
gat_model = GATEncoder(
    in_dim=gat_input_dim,
    hidden_dim=best_config['gat_hidden'],
    out_dim=best_config['gat_hidden'],
    dropout=best_config['gat_dropout'],
    alpha=best_config['gat_alpha']
).to(device)

final_layer = nn.Linear(best_config['gat_hidden'], 1).to(device)

optimizer = torch.optim.Adam(
    list(lstm_model.parameters()) + list(gat_model.parameters()) +
↳list(final_layer.parameters()),
    lr=best_config['learning_rate']
)
```

## 5 Training Loop

```
[44]: epochs = 40
batch_size = best_config['batch_size']
seq_len = 30

date2idx = {d: i for i, d in enumerate(dates)}
train_indices = [date2idx[d] for d in train_dates]

for epoch in range(epochs):
    total_loss = 0.0
    train_points = train_indices[seq_len:-1]
    random.shuffle(train_points)
    num_batches = len(train_points) // batch_size
    if len(train_points) % batch_size != 0:
        num_batches += 1

    progress_bar = tqdm(range(num_batches), desc=f"Epoch {epoch+1}/{epochs}")
```

```

for batch_idx in progress_bar:
    start_idx = batch_idx * batch_size
    end_idx = min(start_idx + batch_size, len(train_points))
    batch_indices = train_points[start_idx:end_idx]
    actual_batch_size = len(batch_indices)

    stock_embeddings_batch = []
    future_returns_batch = []
    past_returns_batch = []

    for t in batch_indices:
        stock_embeddings = []
        future_returns = []

        skip_flag = False
        for stock in tickers:
            df = all_data[stock]
            if t - seq_len < 0 or t + 1 >= len(df):
                skip_flag = True
                break
            seq = df[features].iloc[t-seq_len:t]
            seq = scalers[stock].transform(seq)

            x = torch.tensor(seq, dtype=torch.float32, device=device).
↪unsqueeze(0)
            with torch.no_grad():
                embedding = lstm_model(x).squeeze(0)
                stock_embeddings.append(embedding)
                future_returns.append(df['log_return'].iloc[t+1])
            if skip_flag: continue

        stock_embeddings_batch.append(torch.stack(stock_embeddings))
        future_returns_batch.append(future_returns)
        window = min(20, t)
        past_returns = np.array([
            all_data[stock]['log_return'].iloc[t-window:t].values
            for stock in tickers
        ])
        past_returns_batch.append(past_returns)

    if len(stock_embeddings_batch) == 0: continue

    x_t_batch = torch.stack(stock_embeddings_batch).to(device)
    future_returns_tensor = torch.tensor(future_returns_batch, dtype=torch.
↪float32, device=device)

    batch_loss = 0.0

```

```

    for i in range(x_t_batch.shape[0]):
        x_t = x_t_batch[i]
        updated = gat_model(x_t, edge_index.to(device), edge_attr.
→to(device))
        raw_scores = final_layer(updated).squeeze()
        weights = torch.tanh(raw_scores)
        norm_weights = weights / (weights.sum() + 1e-8)
        cov_matrix = torch.tensor(
            np.cov(past_returns_batch[i]) + 1e-6*np.eye(len(tickers)),
            dtype=torch.float32, device=device
        )
        expected_return = torch.dot(norm_weights, future_returns_tensor[i])
        risk = norm_weights @ cov_matrix @ norm_weights
        loss = -expected_return + 0.1 * risk
        batch_loss += loss

    avg_batch_loss = batch_loss / x_t_batch.shape[0]
    optimizer.zero_grad()
    avg_batch_loss.backward()
    optimizer.step()
    total_loss += avg_batch_loss.item()
    progress_bar.set_postfix(loss=avg_batch_loss.item())

    avg_epoch_loss = total_loss / num_batches
    print(f" Epoch {epoch+1} | Avg Loss: {avg_epoch_loss:.6f}")

print("Training Finished.")

```

```

Epoch 1/40: 100%|      | 12/12 [00:11<00:00,  1.07it/s, loss=-0.00329]
    Epoch 1 | Avg Loss: 0.001687
Epoch 2/40: 100%|      | 12/12 [00:10<00:00,  1.12it/s, loss=-4.51e-5]
    Epoch 2 | Avg Loss: -0.000719
Epoch 3/40: 100%|      | 12/12 [00:10<00:00,  1.11it/s, loss=-0.0031]
    Epoch 3 | Avg Loss: -0.000825
Epoch 4/40: 100%|      | 12/12 [00:10<00:00,  1.13it/s, loss=0.000439]
    Epoch 4 | Avg Loss: -0.000707
Epoch 5/40: 100%|      | 12/12 [00:11<00:00,  1.04it/s, loss=0.000397]
    Epoch 5 | Avg Loss: -0.000728
Epoch 6/40: 100%|      | 12/12 [00:10<00:00,  1.10it/s, loss=0.00121]
    Epoch 6 | Avg Loss: -0.000666
Epoch 7/40: 100%|      | 12/12 [00:10<00:00,  1.11it/s, loss=-0.0029]

```

Epoch 7 | Avg Loss: -0.000799

Epoch 8/40: 100%| | 12/12 [00:10<00:00, 1.11it/s, loss=-0.0002]

Epoch 8 | Avg Loss: -0.000728

Epoch 9/40: 100%| | 12/12 [00:10<00:00, 1.10it/s, loss=-0.00135]

Epoch 9 | Avg Loss: -0.000759

Epoch 10/40: 100%| | 12/12 [00:10<00:00, 1.12it/s, loss=-0.00228]

Epoch 10 | Avg Loss: -0.000821

Epoch 11/40: 100%| | 12/12 [00:10<00:00, 1.09it/s, loss=-0.000196]

Epoch 11 | Avg Loss: -0.000730

Epoch 12/40: 100%| | 12/12 [00:12<00:00, 1.06s/it, loss=-0.00161]

Epoch 12 | Avg Loss: -0.000776

Epoch 13/40: 100%| | 12/12 [00:14<00:00, 1.19s/it, loss=-0.000331]

Epoch 13 | Avg Loss: -0.000697

Epoch 14/40: 100%| | 12/12 [00:11<00:00, 1.08it/s, loss=0.0047]

Epoch 14 | Avg Loss: -0.000542

Epoch 15/40: 100%| | 12/12 [00:10<00:00, 1.10it/s, loss=0.000768]

Epoch 15 | Avg Loss: -0.000679

Epoch 16/40: 100%| | 12/12 [00:11<00:00, 1.08it/s, loss=-0.00326]

Epoch 16 | Avg Loss: -0.000843

Epoch 17/40: 100%| | 12/12 [00:10<00:00, 1.09it/s, loss=-0.00102]

Epoch 17 | Avg Loss: -0.000793

Epoch 18/40: 100%| | 12/12 [00:10<00:00, 1.11it/s, loss=-0.000424]

Epoch 18 | Avg Loss: -0.000733

Epoch 19/40: 100%| | 12/12 [00:10<00:00, 1.10it/s, loss=-0.00156]

Epoch 19 | Avg Loss: -0.000764

Epoch 20/40: 100%| | 12/12 [00:10<00:00, 1.11it/s, loss=0.000941]

Epoch 20 | Avg Loss: -0.000675

Epoch 21/40: 100%| | 12/12 [00:11<00:00, 1.05it/s, loss=-0.00177]

Epoch 21 | Avg Loss: -0.000734

Epoch 22/40: 100%| | 12/12 [00:10<00:00, 1.12it/s, loss=-0.000407]

Epoch 22 | Avg Loss: -0.000734

Epoch 23/40: 100%| | 12/12 [00:10<00:00, 1.16it/s, loss=-0.00363]



Epoch 23 | Avg Loss: -0.000861

Epoch 24/40: 100%| | 12/12 [00:10<00:00, 1.17it/s, loss=-0.0018]

Epoch 24 | Avg Loss: -0.000767

Epoch 25/40: 100%| | 12/12 [00:11<00:00, 1.08it/s, loss=-0.00166]

Epoch 25 | Avg Loss: -0.000768

Epoch 26/40: 100%| | 12/12 [00:13<00:00, 1.09s/it, loss=3.07e-6]

Epoch 26 | Avg Loss: -0.000727

Epoch 27/40: 100%| | 12/12 [00:11<00:00, 1.02it/s, loss=0.00158]

Epoch 27 | Avg Loss: -0.000663

Epoch 28/40: 100%| | 12/12 [00:10<00:00, 1.09it/s, loss=-0.000742]

Epoch 28 | Avg Loss: -0.000770

Epoch 29/40: 100%| | 12/12 [00:10<00:00, 1.11it/s, loss=0.00237]

Epoch 29 | Avg Loss: -0.000640

Epoch 30/40: 100%| | 12/12 [00:10<00:00, 1.13it/s, loss=-0.00121]

Epoch 30 | Avg Loss: -0.000748

Epoch 31/40: 100%| | 12/12 [00:10<00:00, 1.17it/s, loss=-0.00534]

Epoch 31 | Avg Loss: -0.000884

Epoch 32/40: 100%| | 12/12 [00:10<00:00, 1.16it/s, loss=-0.00495]

Epoch 32 | Avg Loss: -0.000859

Epoch 33/40: 100%| | 12/12 [00:10<00:00, 1.17it/s, loss=-3.71e-5]

Epoch 33 | Avg Loss: -0.000701

Epoch 34/40: 100%| | 12/12 [00:10<00:00, 1.16it/s, loss=0.00056]

Epoch 34 | Avg Loss: -0.000704

Epoch 35/40: 100%| | 12/12 [00:10<00:00, 1.18it/s, loss=0.00639]

Epoch 35 | Avg Loss: -0.000515

Epoch 36/40: 100%| | 12/12 [00:10<00:00, 1.16it/s, loss=-0.00234]

Epoch 36 | Avg Loss: -0.000802

Epoch 37/40: 100%| | 12/12 [00:10<00:00, 1.17it/s, loss=-0.0025]

Epoch 37 | Avg Loss: -0.000809

Epoch 38/40: 100%| | 12/12 [00:10<00:00, 1.15it/s, loss=0.00114]

Epoch 38 | Avg Loss: -0.000687

Epoch 39/40: 100%| | 12/12 [00:10<00:00, 1.16it/s, loss=-0.000681]

Epoch 39 | Avg Loss: -0.000717  
Epoch 40/40: 100%| | 12/12 [00:10<00:00, 1.09it/s, loss=-0.000527]  
Epoch 40 | Avg Loss: -0.000748  
Training Finished.

```
[45]: # Save
torch.save(lstm_model.state_dict(), "lstm_model.pth")
torch.save(gat_model.state_dict(), "gat_model.pth")
torch.save(final_layer.state_dict(), "final_layer.pth")

# Load
lstm_model.load_state_dict(torch.load("lstm_model.pth"))
gat_model.load_state_dict(torch.load("gat_model.pth"))
final_layer.load_state_dict(torch.load("final_layer.pth"))
lstm_model.eval()
gat_model.eval()
final_layer.eval()
print("Complete saving and loading model.")
```

Complete saving and loading model.

## 6 Start Testing

```
[46]: lstm_model.eval()
gat_model.eval()
final_layer.eval()

test_indices = [date2idx[d] for d in test_dates]
test_points = test_indices[seq_len:-1]

weights_all_days = []
test_dates_list = []

num_batches = len(test_points) // batch_size
if len(test_points) % batch_size != 0:
    num_batches += 1

for batch_idx in tqdm(range(num_batches), desc="Predicting (test set)":
    start_idx = batch_idx * batch_size
    end_idx = min(start_idx + batch_size, len(test_points))
    batch_indices = test_points[start_idx:end_idx]

    stock_embeddings_batch = []

    for t in batch_indices:
```

```

stock_embeddings = []
skip_flag = False
for stock in tickers:
    df = all_data[stock]
    if t - seq_len < 0 or t + 1 >= len(df):
        skip_flag = True
        break
    seq = df[features].iloc[t-seq_len:t]
    seq = scalers[stock].transform(seq)
    x = torch.tensor(seq, dtype=torch.float32, device=device).
↳unsqueeze(0)
    with torch.no_grad():
        embedding = lstm_model(x).squeeze(0)
        stock_embeddings.append(embedding)
    if skip_flag: continue
    stock_embeddings_batch.append(torch.stack(stock_embeddings))
    test_dates_list.append(dates[t])

if len(stock_embeddings_batch) == 0: continue

x_t_batch = torch.stack(stock_embeddings_batch).to(device)

for i in range(x_t_batch.shape[0]):
    x_t = x_t_batch[i]
    with torch.no_grad():
        updated = gat_model(x_t, edge_index.to(device), edge_attr.
↳to(device))
        raw_scores = final_layer(updated).squeeze()
        weights = torch.tanh(raw_scores)
        norm_weights = weights / (weights.sum() + 1e-8)
        weights_all_days.append(norm_weights.cpu().numpy())

# Save predicted weight
weights_df = pd.DataFrame(weights_all_days, columns=tickers)
weights_df['date'] = test_dates_list
weights_df.to_csv("Result/predicted_weights_v1.csv", index=False)
print("File path: Result/predicted_weights_v1.csv")

```

Predicting (test set): 100% | 5/5 [00:03<00:00, 1.41it/s]

File path: Result/predicted\_weights\_v1.csv

```

[ ]: # Build testing DataFrame
test_price_data = pd.DataFrame({'date': test_dates_list})
for stock in tickers:
    prices = []

```

```

for d in test_dates_list:
    idx = date2idx[d]
    prices.append(all_data[stock].iloc[idx]['close'])
test_price_data[stock] = prices

# Daily return
returns = test_price_data[tickers].pct_change().dropna().reset_index(drop=True)
weights = weights_df[tickers].iloc[:-1].reset_index(drop=True)

# Portfolio return and equal-weight return
portfolio_returns = (returns.values * weights.values).sum(axis=1)
equal_weights = np.ones(len(tickers)) / len(tickers)
equal_returns = (returns.values * equal_weights).sum(axis=1)

# Cumulative Return
cumulative_portfolio = (1 + portfolio_returns).cumprod()
cumulative_equal = (1 + equal_returns).cumprod()


# Visualized
import matplotlib.pyplot as plt
import matplotlib.dates as mdates

plot_dates = pd.to_datetime(test_price_data['date'].iloc[1:])

plt.figure(figsize=(12, 8))
plt.plot(plot_dates, cumulative_portfolio, label='LSTM-GAT Portfolio (v1)',
         linewidth=2)
plt.title('Out-of-Sample Portfolio Performance', fontsize=16)
plt.xlabel('Date', fontsize=14)
plt.ylabel('Cumulative Return', fontsize=14)
plt.legend(fontsize=12)
plt.grid(True, alpha=0.3)
plt.gca().xaxis.set_major_locator(mdates.AutoDateLocator())
plt.gca().xaxis.set_major_formatter(mdates.DateFormatter('%Y-%m-%d'))
plt.gcf().autofmt_xdate()
plt.tight_layout()
plt.savefig('Result/V1out_of_sample_performance.png', dpi=300)
plt.show()

print("Testing set cumulative return graph finished.")

```



Testing set cumulative return graph finished.

```
[48]: # Result DataFrame
result_df = pd.DataFrame({
    'date': plot_dates,
    'lstm_gat_return': portfolio_returns,
    'lstm_gat_cum_return': cumulative_portfolio,
    'equal_weight_return': equal_returns,
    'equal_weight_cum_return': cumulative_equal
})

result_df.to_csv('Result/portfolio_returns_v1.csv', index=False)
print("File path: Result/portfolio_returns_v1.csv")
```


File path: Result/portfolio\_returns\_v1.csv

```
[ ]: capm_df = (
    pd.read_csv("Result/capm_cumulative_returns.csv",
        parse_dates=["date"])
    .rename(columns={"cumulative_return": "capm_cumulative"})
)

# If you only want the dates that appear in your LSTM test window:
capm_df = capm_df[capm_df["date"].between(
```

```
test_price_data['date'].iloc[1],  
test_price_data['date'].iloc[-1]])
```

```
[ ]: # Plot LSTM-GAT vs Equal-Weight vs CAPM Max-Sharpe  
plt.figure(figsize=(12, 8))  
  
# already in your code  
plt.plot(plot_dates, cumulative_portfolio,  
         label='LSTM-GAT Portfolio', linewidth=2)  
plt.plot(plot_dates, cumulative_equal,  
         label='Equal-Weight Portfolio', linestyle='--')  
  
# NEW: CAPM curve (dotted line)  
plt.plot(capm_df['date'], capm_df['capm_cumulative'],  
         label='CAPM Max-Sharpe Portfolio', linestyle=':')  
  
plt.title('Out-of-Sample Portfolio Performance (v1)', fontsize=16)  
plt.xlabel('Date', fontsize=14)  
plt.ylabel('Cumulative Return', fontsize=14)  
plt.legend(fontsize=12)  
plt.grid(True, alpha=0.3)  
plt.gca().xaxis.set_major_locator(mdates.AutoDateLocator())  
plt.gca().xaxis.set_major_formatter(mdates.DateFormatter('%Y-%m-%d'))  
plt.gcf().autofmt_xdate()  
plt.tight_layout()  
  
plt.savefig('Result/V1_out_of_sample_performance_with_capm.png', dpi=300)  
plt.show()
```



```
[51]: import matplotlib.pyplot as plt
import matplotlib.dates as mdates # For formatting date ticks


plt.figure(figsize=(14, 7))
weights_df['date'] = pd.to_datetime(weights_df['date'])

# Plot portfolio weights
for stock in tickers:
    plt.plot(weights_df['date'], weights_df[stock], label=stock)

plt.gca().xaxis.set_major_locator(mdates.AutoDateLocator())
plt.gca().xaxis.set_major_formatter(mdates.DateFormatter('%Y-%m-%d'))
plt.xticks(rotation=45)

plt.xlabel('Date')
plt.ylabel('Portfolio Weight')
plt.title('Predicted Portfolio Weights on Test Set (v1)')
plt.legend()

plt.tight_layout()
plt.show()
```



## 7 CAPM-MVO Result

```
[52]: capm_weights_df = (
    pd.read_csv("Result/capm_daily_weights.csv", parse_dates=["date"])
    .sort_values("date")
)

TICKERS = ["AAPL", "AMZN", "BA", "COST", "JNJ", "NVDA", "TMO", "TSLA", "VLO"]

# Plot weights
import matplotlib.pyplot as plt
import matplotlib.dates as mdates

plt.figure(figsize=(12, 7))


for stk in TICKERS:
    plt.plot(capm_weights_df["date"],
             capm_weights_df[stk],
             label=stk, linewidth=1.8)

plt.title("CAPM Portfolio Weights (Test Period)", fontsize=16)
plt.xlabel("Date", fontsize=14)
plt.ylabel("Weight", fontsize=14)
plt.axhline(0, color="black", linewidth=0.8)
plt.gca().xaxis.set_major_locator(mdates.AutoDateLocator())
plt.gca().xaxis.set_major_formatter(mdates.DateFormatter('%Y-%m-%d'))
```



```
plt.gcf().autofmt_xdate()

plt.grid(alpha=0.3)
plt.legend(ncol=3, fontsize=10, frameon=False)
plt.tight_layout()
plt.savefig("Result/capm_weight_paths.png", dpi=300)
plt.show()
```



## 8 Evaluation

```
[ ]: import numpy as np
import pandas as pd

def calculate_metrics(returns, var_conf_level=0.95):
    returns = pd.Series(returns)
    cumulative = (1 + returns).cumprod()
    total_return = cumulative.iloc[-1] - 1
    annualized_return = (1 + total_return) ** (252 / len(returns)) - 1
    volatility = returns.std() * np.sqrt(252)
    sharpe_ratio = annualized_return / volatility if volatility > 0 else 0
    max_drawdown = (cumulative / cumulative.cummax() - 1).min()

    # Historical Value at Risk (e.g., 5% worst return)
    var_percentile = 100 * (1 - var_conf_level)
    value_at_risk = -np.percentile(returns, var_percentile)
```

```

    return total_return, annualized_return, volatility, sharpe_ratio,
    ↪max_drawdown, value_at_risk

```

```

port_metrics = calculate_metrics(portfolio_returns)
equal_metrics = calculate_metrics(equal_returns)
capm_df = pd.read_csv('Result/capm_daily_returns.csv')
capm_returns = capm_df['daily_return'].values[1:]
capm_metrics = calculate_metrics(np.exp(capm_returns) - 1)

```

```

[62]: print("\n" + "="*80)
print("Out-of-Sample Performance Comparison (Test Period)")
print("="*80)
print(f"{'Metric':<20}{'LSTM-GAT (v1)':>20}{'Equal-Weight':>20}{'CAPM':>20}")
print(f"{'Total Return':<20}{port_metrics[0]:>20.6%}{equal_metrics[0]:>20.6%}{capm_metrics[0]:>20.6%}")
print(f"{'Annualized Return':<20}{port_metrics[1]:>20.6%}{equal_metrics[1]:>20.6%}{capm_metrics[1]:>20.6%}")
print(f"{'Volatility':<20}{port_metrics[2]:>20.6%}{equal_metrics[2]:>20.6%}{capm_metrics[2]:>20.6%}")
print(f"{'Sharpe Ratio':<20}{port_metrics[3]:>20.6f}{equal_metrics[3]:>20.6f}{capm_metrics[3]:>20.6f}")
print(f"{'VaR (95%)':<20}{port_metrics[5]:>20.6%}{equal_metrics[5]:>20.6%}{capm_metrics[5]:>20.6%}")
print(f"{'Max Drawdown':<20}{port_metrics[4]:>20.6%}{equal_metrics[4]:>20.6%}{capm_metrics[4]:>20.6%}")
print("="*80)

```

```

=====
Out-of-Sample Performance Comparison (Test Period)
=====

```

Metric	LSTM-GAT (v1)	Equal-Weight	CAPM
Total Return	29.937403%	26.556317%	24.340720%
Annualized Return	24.605357%	21.876063%	20.607402%
Volatility	25.572940%	24.813493%	22.234806%
Sharpe Ratio	0.962164	0.881620	0.926808
VaR (95%)	2.558731%	2.514239%	2.022031%
Max Drawdown	-23.074754%	-22.604449%	-21.032717%

```

=====

```